

What four years of testing actually proved

Evidence notes for discretionary trading

Every claim in this document is marked with what the evidence says, how strong it is, and **how many observations it rests on**. Nothing here is ICT orthodoxy repeated back — where a belief was tested and failed, it is in the DON'T section with the number that killed it.

The evidence base

Source	Period	Observations
Main backtest (the algorithm)	2022–2025, 3 pairs	736 trades · 355 in-sample / 381 out
Tick-volume AMD study	2022 & 2024, EU+GU	275 trades
Pure-price draw cascade	2022–2025, 3 pairs	3,637 sweep events
Triple-raid study	2022–2025, H1	3,045 raids (1,541 short / 1,504 long)
Gap-as-magnet study	2022–2025, W/D/H4	62,797 gaps
IFVG standalone backtest	2022–2025, 3 pairs	~10,000 simulated trades
Gap-race study	2022–2025, daily	3,157 bars · 774 resolved races
Levers actually tested in the engine	—	~25 (5 shipped, 20 rejected)

The single most important framing. Most of what sounds true in ICT is *descriptively* true and carries no directional information. Price does go to the gap. The gap does hold. Both are real and both are useless on their own, because price does the same thing at any nearby level. Everything below separates the two.

Baseline the algorithm produced: 736 trades, win rate 43.9%, profit factor 4.01, worst drawdown –13.24%. R1,000 grew to R132,020 withdrawn plus R8,555 working.

Part 1 — The DOs

Held in **both** halves of the four years, at a similar size, on enough observations to mean something.

CONFIRMED

1. The AMD cycle is real, and you can see it in participation

Tick volume measured against each trade's own pre-accumulation baseline:

Phase	Volume vs baseline	Observations
Accumulation (the coil)	0.66× — dries to two-thirds	275 trades
Manipulation (the stop run)	1.26× median, 1.74× mean	275 trades
Distribution	1.22×	275 trades
At the PD array (the fill)	1.28× — reacts hardest	275 trades

Held on both pairs — GBPUSD spikes harder than EURUSD on the sweep (1.93× vs 1.63×). **When a range goes quiet and then a sweep arrives on a violent bar, that is the real thing, not a fakeout.**

But volume does NOT separate winners from losers. Tested directly on the same 275 trades: the fingerprints are the same. The small 2022 edge flips sign in 2024. Absorption at the sweep is 50% win / 50% lose (2022) and 48% / 48% (2024). Use volume to confirm you are *in* an AMD cycle — never to judge whether *this* one pays.

CONFIRMED

2. A sweep that runs prior-day high/low is your better setup

What the sweep ran	2022 win rate	2024 win rate	Observations
Prior-day high / low	55%	53%	115 of 810 trades
Everything else (baseline)	45.3%	45.8%	695 of 810

+8 to +10 points in **both** years. Shipped as a 1.25× size increase; it fired on 122 trades over four years and lifted equity **+57%** with drawdown held to the decimal (−12.95% before and after). It is the one setup-quality signal in the whole project that earned more size on merit.

For scale, what sweeps actually run across 810 trades: equal highs/lows **397**, prior-day **115**, prior-week **8**, value area **9**.

CONFIRMED

3. The H4 turtle soup is a genuine timing confirmation

Price wicks beyond the prior two-bar H4 range, then closes back inside — the Judas swing one timeframe up. The manipulation has already fired.

Higher-timeframe sweep	Trades	Win rate	PF 2022-23	PF 2024-25
H4 turtle soup	137	50.4%	3.40	3.80
Daily turtle soup	234	46.6%	2.88	3.02
No sweep	427	45.4%	3.21	8.11

The H4 one is the good one; the daily version is weaker. Shipped as a 1.25× size increase on the H4 bucket only — 122 trades sized, equity **+48%**, drawdown unchanged at −12.95%.

CONFIRMED

4. The golden rule — it is about WHICH PAIR, not just direction

	2022-23 WR / PF	2024-25 WR / PF
SELL GBPUSD · BUY EURUSD	47.5% / 4.07	50.0% / 4.39
Against the rule	39.7% / 2.84	42.2% / 4.06

+8 points in both halves, same size. GBP distributes harder on the downside, EUR on the upside. **Short the weaker, buy the stronger.** Shipped as a 1.25× size increase.

Related and consistent: an order-block entry on GBPUSD ran WR 60.0% (2022-23) and 64.3% (2024-25) — the standout combination. Order blocks beat gaps on profit factor in both halves (3.72 / 5.18 vs 3.31 / 3.54), though gaps carry the volume at 484 trades.

The fuller four-quadrant version uses EURGBP to pick the pair, and the counters show **EURGBP is the heavier filter** — it rejected 97 setups where the dollar rejected 59. The cross does more work than the dollar. That version was never validated out of sample, so treat the quadrant as a framework, not a proven edge. The simple rule above IS proven.

CONFIRMED

5. Three independent confluences at your target — and not more

Each target scored by how many *different* families agree within 8 pips: fib, gap, order block, equal highs/lows, round number, prior-day, prior-week, intermediate swings.

Confluences	Trades	Profit factor
1	13	0.02
2	40	0.20
3	225	6.70
4	325	5.73
5	190	3.94
6	18	43.28 (tiny sample)

The cliff between 2 and 3 is the finding — a 33× jump in profit factor across 265 trades. One source can be found near any price; that is noise. Three different frameworks pointing at the same level means multiple desks are aiming there.

Note it does **not** keep improving above 4. More confirmation past that point adds nothing — and costs you entry price.

CONFIRMED

6. Stop one structural tier beyond the obvious swing

Change	Equity	Profit factor	Worst drawdown
Stop at the nearest short-term swing	R60.18M	5.03	−12.95%
Stop beyond the intermediate swing	R71.19M (+18.3%)	5.12	−12.95%

The short-term swing is exactly what a minor liquidity run sweeps. That is what it is there for. A stop at the obvious recent low is parked in the pool. One tier up survives the sweep — and costs a little more on the trades where you were simply wrong, which is the right trade-off.

CONFIRMED

7. The 3-day pool is the dependable draw. The 30/60-day is not.

Measured on raw price across **3,637 prior-day sweep events**, independent of any strategy. After a sweep, which pool does price reach within two days:

Next pool	2022-23	2024-25	Read
3-day high / low	58%	61%	the dependable objective
Weekly	25%	27%	occasional
30-day	21%	20%	extension only
60-day	15%	13%	extension only

All three pairs agree (roughly 60 / 20 / 14). **Target the 3-day pool.** The 30 and 60-day pools are reached about one time in five — they are extension targets, not objectives. This is the arithmetic behind DON'T #2 in Part 3.

CONFIRMED

8. Trail as it runs on a long trade

Breakeven at +10, lock +10 at +20, then every 20 pips of further progress lock 10 behind:

	Equity	Profit factor	Worst drawdown
Before	R294.8M	4.47	−12.95%
With progressive trailing	R400.7M (+36%)	4.47	−13.01%

The largest single improvement in the project after position sizing, for +0.06% of drawdown. **A trade that reaches +35 and reverses should not be a full loss.**

CONFIRMED

9. Monday and Tuesday of NFP week

	2022-23 WR / PF	2024-25 WR / PF
Mon–Tue of NFP week	62.5% / 4.66	52.1% / 6.12
All other days	41.2% / 3.25	43.3% / 4.03

The only one of **seven** contextual factors that passed both halves. The tight pre-NFP accumulation makes for a clean sweep. Wednesday to Friday of that week is already a known low-probability window.

Sample counts per bucket are not broken out in the record — treat this as the weakest of the nine DOs on evidence, even though it passed.

Part 2 — The DON'Ts

Things that sound true, that most people believe, and that measure to nothing once you ask the right control question: **what would a meaningless level do?**

MEASURED NULL

1. “Price is drawn to the fair value gap”

True — and worthless as a direction. Across **62,797 gaps**, compared against a band of identical width, the same distance away, on the *opposite* side of price:

Timeframe	Gap reached	Mirror band reached	Difference
Weekly	69.7% / 78.8%	71.2% / 80.3%	−1.5 / −1.5
Daily	90.2% / 87.2%	91.0% / 86.4%	−0.7 / +0.8
H4	95.2% / 95.3%	95.0% / 95.3%	+0.2 / +0.0

Six comparisons, **none differing by more than 0.8 points**, two negative. On completely random data both read 93.5%. Adding a structure shift toward the gap made it *worse* (85.1% / 86.0% against its own control's 92.6% / 96.0%).

The gap being reached tells you nothing about which way to trade. Holding a bias toward a gap for days adds risk and time without adding expectancy. If you catch yourself saying “price has to come back and fill that gap” — it does, and so does the level the same distance in the other direction.

MEASURED NULL

2. “It held at the gap, so it's a support zone”

Timeframe	Gap respected	Mirror band respected	Difference
Weekly	69.6% / 69.2%	63.8% / 73.6%	+5.7 / −4.4
Daily	70.2% / 66.9%	60.5% / 63.4%	+9.7 / +3.5
H4 (n ≈ 2,000 per half)	64.2% / 66.4%	63.7% / 64.3%	+0.6 / +2.1

The one cell that looked real — daily, +9.7 points — collapsed to +3.5 in the second half. And it **lost money**: after a “respected” daily gap the favourable move averaged 86 pips against 98 pips adverse. A higher reversal rate that loses money is not a signal.

On H4 the sample is about 2,000 per half, where the margin of error is only 1.5 points — so the flat reading there is a real measurement, not a small-sample shrug.

MEASURED NULL

3. “A broken gap inverts and becomes resistance” (the IFVG)

Tested as a standalone entry with a realistic 10-pip structural stop and costs, across three pairs and four years — **roughly 10,000 simulated trades**:

Timeframe	Trades (IS / OOS)	PF 2022-23	PF 2024-25	Win rate
M15	4,022 / 4,543	0.78	0.73	33% / 32%
H1	1,243 / 1,122	0.83	0.85	33% / 34%
H4	378 / 292	0.97	0.89	37% / 35%
Daily	89 / 78	0.86	0.67	34% / 28%

Below 1.0 on every timeframe, in both halves, on two different target methods. Median result −1.11R — most trades stop out. The reason is arithmetic: a 2R target off a 10-pip stop needs better than 33% to break even before spread, and the inversion lands right at 33%.

Tested a second, independent way (does the zone reject price on a re-test): hold rates look strong at 63–78%, but the control does 61–75%. Every difference under one standard error.

The full-body-close inversion is a real market-structure event that carries no measurable edge on its own. It can still be useful as *context* inside a bigger setup — that is what the semi-auto model uses it for. It is not a trade by itself.

MEASURED NULL

4. “Both pairs took highs while the dollar took lows — triple confirmation”

	Triple raid	Single-pair raid	Observations
Short side	52.3% / 50.5%	50.6% / 49.1%	1,541 trios
Long side	48.7% / 51.7%	47.8% / 52.3%	1,504 trios

Every bucket, both directions, both halves, lands between 47.5% and 52.3%. The margin of error at that sample size is ± 1.9 points. A lone sweep with no alignment does exactly the same thing.

3,045 of these in four years — about 1.5 per day. It is not a rare institutional fingerprint; it is what the dollar complex does mechanically most days, because the pairs are inverse to the dollar. **You are counting one event three times.**

Adding the dollar/cross quadrant on top made it *worse* in all four cells (−4.8, −3.2, −0.5, −2.3 in the first half).

MEASURED NULL

5. “The daily open is a key level”

Tested against a deliberately meaningless level — a quarter of yesterday's range from the open, side alternating by date. Lift over that placebo:

Pair	2022-23	2024-25	Read
EURUSD (ran the low)	+2.1	−4.4	sign flips
GBPUSD (ran the high)	−3.2	+3.7	sign flips
NZDUSD (ran the low)	−5.9	+0.2	sign flips

The daily open is not special. Same verdict for the session open.

MEASURED NULL

6. “More confirmations means a better trade”

Counting how many opening references confirmed (0, 1 or 2):

Confirmations	2022-23 (trades / WR)	2024-25 (trades / WR)
Both references (+2)	123 / 48.0%	100 / 41.0%
One reference (+1)	188 / 41.5%	246 / 45.1%
None	44 / 36.4%	39 / 48.7%

Perfectly ordered downward in the first half and perfectly ordered upward in the second, on comparable samples. Two mirror images is the clearest possible picture of noise. Combined with the confluence table flattening above 4, the message is the same: waiting for a fourth and fifth confirmation buys you nothing and costs you entry price.

MEASURED NULL

7. “If my stop was hit, I was early — the move came later”

Your own standing claim, and worth testing properly. From the bar after each losing trade closed: which comes first, the level we aimed at, or a mirror level the same distance the *other* way?

Half	Ours first	Mirror first	Both	Neither	Ours %
2022-23 (losses)	57	50	1	48	53.3%
2024-25 (losses)	71	74	0	40	49.0%

A coin flip in both halves (107 and 145 resolved races, margin of error ± 4.5 points). The losses were wrong-way trades, not early ones — so no stop placement, entry timing or extra patience recovers them. Worth sitting with, because “I was right but early” is the most comfortable story a trader tells after a loss, and here it is measurably false.

8. Day of week, seasonal lean, and “narrative” generally

Seven contextual factors, each scored on all 740 trades. One passed (NFP Mon–Tue). One was mildly **inverted** — the trades where the prior session's sweep agreed with the direction were *worse* in both halves. One never fired at all. The remaining four flipped sign between halves.

The aggregate score was not even ordinal: in the first half, score 1 → PF 0.93, 2 → 4.17, 3 → 2.94, 4 → 3.82, 5 → 15.01.

Part 3 — What went wrong when we added things

Five patterns, each tested in the engine with real money at stake, each one costly. **These translate most directly to discretionary decisions.**

LOST MONEY

1. Holding past your target — three different ways, all destroyed money

Attempt	Equity	Change
Baseline	R284.9M	—
Trail instead of exiting on a target wick (v1)	bankrupt	–100.6% drawdown
Same, with a breakeven guard (v2)	R137.8M	–52%
Same, best implementation (v3)	R144.8M	–49%
Take half at target, run the rest to the next draw	R70M	–75%

Every single exit modification lost. The first version was so bad it bankrupted the test account.

Why: your target IS the end of the AMD cycle. When price delivers to the institutional draw, distribution has happened. Holding past it means sitting through the consolidation and reversal where you have no edge. **Take the full position off at the draw.** The 100–200 pip continuation you sometimes see afterwards is real — and it is a new trade with a new entry, not a reason to hold.

LOST MONEY

2. Reaching for the further target destroyed money

	Worst drawdown 2022-23	Worst drawdown 2024-25	Win rate
Nearest qualifying target	–13.01%	–13.94%	46.0% / 45.9%
Prefer the distant pool	–22.57%	–22.60%	41.8% / 42.5%

Drawdown nearly doubled in both halves. **The near target is the draw for THIS cycle; the far one belongs to the NEXT cycle.** Aiming at the far one means staying through the intermediate consolidation where price pauses or reverses. The 58%-vs-21% reach rates in DO #7 are the arithmetic behind it.

LOST MONEY

3. Trading INTO the pool destroyed money — four attempts

	Worst drawdown	Profit factor	Win rate
Fade the sweep (the shipped model)	–12.95%	3.49	45.3%
Trade toward the untouched pool	–16.73%	3.30	44.6%

72 continuation entries added; they won less often and clustered their losses, breaking the risk limit. Four separate attempts at “continuation instead of reversal” failed the same way.

The pool is where the cycle ENDS. Riding into it means buying exactly where smart money is selling into the resting liquidity. **Prior-day highs and lows are correct as targets. They are wrong as entries.**

LOST MONEY

4. Filtering out “bad” setups destroyed money — the least intuitive one

Filter tried	Bucket quality	Cost
Cut a clearly losing session phase	PF 0.13 / 0.16	–R31.1M of compounding
Merely shrink the weaker NY echo	WR 46.9%	–R5.8M, lost in all 3 tests
Skip reversals fighting an opposing gap	passed both halves	–20.15% drawdown on the full run
Exempt 6 weak trades from a gate	6 trades	–R10.3M (–17%)

A bucket with a profit factor of **0.13** — catastrophic in isolation — cost **R31 million** of compounded growth when removed. A gentler version that merely shrank a weak bucket lost equity in all three test periods and improved drawdown in none.

Why: a setup that loses on simple arithmetic can still be positive to your compounding, because its occasional wins land at points where they compound forward, and the slot it occupies would otherwise go to a lower-quality trade. **Be very careful about deciding a setup type is bad and cutting it.** The obvious filter is usually wrong — four separate attempts, all costly. This is the single most repeated failure in the project.

LOST MONEY

5. Sizing up broadly is leverage, not selection

Size increase applied to...	Trades	Profit factor	Equity
69% of the book (505 of 736)	736 → 696	4.01 → 3.58	R140.6k → R137.5k
A gap at its midpoint (shipped)	fired 15×	unchanged	+R825k
Prior-day sweep (shipped)	fired 91×	unchanged	+57%
H4 turtle soup (shipped)	fired 122×	unchanged	+48%

The broad version *lowered* profit factor and *reduced* the trade count — the bigger positions breached the per-trade risk cap and got skipped, 414 times.

Every size increase that worked targeted a small, distinctive subset — 15, 91 or 122 trades out of 736. A size bump is a selection tool. Apply it to most of your trades and you have simply raised your lot size.

Part 4 — How not to fool yourself

Method lessons, and honestly the most transferable part. I made every one of these mistakes during this work.

METHOD

1. An unusually clean result is evidence of a bug before it's evidence of an edge

Five separate studies produced beautiful, publishable findings that were entirely artifacts of my own measurement. In every single case the broken version looked *better* than the truth.

One of them — a target-quality finding I had called the strongest result in the project — turned out to be reading information from the future. When fixed it did not shrink, it **inverted**: the bucket I had called a loser (profit factor 0.84) was actually the best one (5.72). Acting on the broken version would have cost 42% of equity.

The same class of error appeared **seven times** across the project.

METHOD

2. Always ask what a meaningless level would do

This is what killed most of Part 2. “The gap gets reached 95% of the time” sounds decisive until you check that a random band the same distance away gets reached 95% too — and that on pure random data both read 93.5%.

Before you credit a level for what price did, ask what price does at levels generally.

METHOD

3. One good period proves nothing

Three times in this project a perfectly clean, ordered pattern in 2022-23 dissolved — or exactly reversed — in 2024-25. Not degraded. Mirror image.

If you find something that works, you have not found something that works. You have found something that *worked in the period you looked at*.

METHOD

4. Don't let a confirmation be the thing you're trying to predict

Four times I measured “structure shifted my way” over the same window as “price moved my way” — which are the same statement. The bucket scored 100% by construction. On one check, 62,797 gaps produced **zero** exceptions to the circularity.

If your confirmation only becomes visible after the move has already happened, it isn't a confirmation.

METHOD

5. Know how much data your question needs before you cut it

The final study needed roughly **four times the entire four-year dataset** to answer its question: 774 resolved observations against about 1,580 required. Every refinement made it worse by subdividing further.

With 40 trades a quarter you will never resolve a 5-point win-rate difference. Don't try. Spend your attention where the effect is large enough to see.

Part 5 — Your edge versus the algorithm's

There is direct evidence in this project that a human can beat the machine on exactly this material.

The Market Maker model — dealing range, sweep of external liquidity, entry at the inverted gap — was tested autonomously across four years:

Measure	Result
Setups produced	1,091–1,158 over 4 years (~5.5 per week)
Money profit factor	2.98 — genuinely profitable
Win rate	21%, and uniform across every tag
Shape of the edge	845 losers at ~-4 pips, 230 winners at ~+15, tail to +114
Worst drawdown	-23.4%, improved only to -21.2% by every lever tried

It was not shipped for exactly one reason, quoted from the record:

“Good timeframe cascades with a profit factor above 1 in both halves: **None**. Win rate 21% is uniform across every tag — cascade timeframe, pattern, pair, period. No robust subset to filter to. The edge is in a fat tail of rare large winners, **and the winners are unpredictable in advance**.

Unpredictable in advance **by a machine, from the tags it can measure**. That is precisely the gap a discretionary trader fills. You are reading context the classifier cannot encode — how the sweep felt, what the dollar was doing into it, whether the session had already shown its hand. The machine had to reject that setup. You don't.

So yes — you may well make more than the algorithm on this material. That is a reasonable expectation, not wishful thinking.

What the algorithm has that you don't

- It takes the boring setup at 3:07am without being bored.
- It never revenge trades after two losses.
- It never skips the seventh small setup because the last six annoyed it.
- It stops at -6% on the day, every day, without negotiating.
- It sizes identically on the setup that feels great and the one that doesn't — and the record above says that feeling is mostly not informative.

What you have that it doesn't: **context**. The thing that made the model unfilterable.

The honest split. Your edge is in *selection* — picking which of the day's setups is the real one. The machine's edge is in *discipline* — executing without mood. The worst version of you is one who has the context read right and gives it back through size, revenge, or skipping the boring ones.

Keep the circuit breakers as personal rules even trading by hand

Two losses, stop for the day. And **no fresh trade in the same dollar direction as one that just stopped out** — EURUSD, GBPUSD and NZDUSD are one dollar bet. Blocking concurrent same-direction exposure is what lifted that model's profit factor from **2.47 to 2.98**. That one is measured, not advice.

The one-page version

DO

	Evidence
Trade the pair the dollar and EURGBP select — sell the weaker, buy the stronger	+8pp both halves
Prefer sweeps that ran yesterday's high or low	55% / 53% vs 45%, 115 trades
Prefer setups where the H4 candle already swept and rejected	50.4% WR, 137 trades
Need three <i>different</i> things agreeing at your target — not more	PF 0.20 → 6.70 at three
Stop one structural tier beyond the obvious swing	+18% equity, drawdown flat
Target the 3-day pool; take the nearest target clearing your R:R	58% / 61% reached
Trail as it runs; take the whole position off at the draw	+36% equity
Mon–Tue of NFP week is a good window	62.5% / 52.1% WR

DON'T

	What killed it
Don't hold past your target hoping for continuation	–49% and –75%
Don't reach for the further target	drawdown –13% → –22%
Don't trade <i>into</i> an untouched pool — that's where the cycle ends	–16.7%, 4 attempts
Don't decide a setup type is bad and cut it	–R31M on a PF 0.13 bucket
Don't wait for a fourth and fifth confirmation	flat above 3; mirror-image halves
Don't credit the gap for price arriving — it arrives everywhere	mirror band: same 90–95%
Don't tell yourself you were early	53.3% / 49.0% — coin flip
Don't take a fresh trade in the same dollar direction as one that just stopped	PF 2.47 → 2.98 when blocked

And the meta-rule: if something looks unusually clean, it is probably a measurement artifact. That held **five times out of six** in this project — and in every case the broken version looked better than the truth.

Sources: every figure traces to a measured result recorded in CLAUDE.md — the running research log for this project.
Baseline algorithm: 736 trades over 2022–2025, win rate 43.9%, profit factor 4.01, worst drawdown –13.24%.